

On some modern developments in generative modelling

ISA Oberseminar

Lukas Trottner

18 November 2025

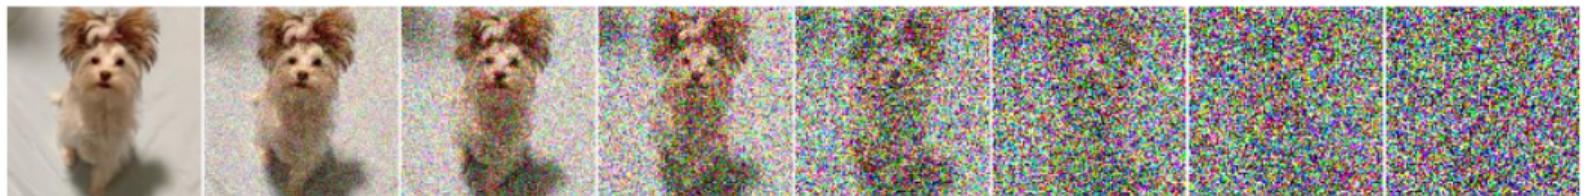
University of Stuttgart



University of Stuttgart
Germany

Motivation:

“Creating noise from data is easy; creating data from noise is generative modeling.”



Source: Song et al. (2021). Score based generative modeling through stochastic differential equations. *ICLR*.

Generative modelling

Setup: identically distributed samples X_1, \dots, X_n with **unknown distribution P** are given

Goal: develop sampling algorithms that do not rely on structural assumptions on P

- ~~ involves (implicitly) **learning the underlying distribution** of a dataset to generate new samples that
 - a) follow approximately the same distribution as the training data;
 - b) should not be drawn from the training data set
- ~~ essential in applications like image synthesis, text generation, data augmentation ...

Noise transformation

Inverse transform sampling: for an \mathbb{R} -valued random variable X with cdf F and $U \sim \mathcal{U}((0, 1))$, we have $F^{-1}(U) \stackrel{d}{=} X$ for the left-inverse F^{-1} of F

- we don't know F , but are only given samples $X_1, \dots, X_n \stackrel{d}{=} X$
 - naïve approach: replace F by empirical cdf $\hat{F}(x) = \frac{\#\{X_i : X_i \leq x\}}{n}$ and set $\hat{X} = \hat{F}^{-1}(U)$ for an independent $U \sim \mathcal{U}((0, 1))$
 - if $X_{(1)}, \dots, X_{(n)}$ is an increasing ordering of the data set and $U \in [k/n, (k+1)/n]$, then $\hat{X} = X_{(k)}$
- ⇒ algorithm learns the empirical distribution $\mathbb{P}_n = \frac{1}{n} \sum_{i=1}^n \delta_{X_i}$ ⇒ overfitting/no creativity"

Noise transformation

Inverse transform sampling: for an \mathbb{R} -valued random variable X with cdf F and $U \sim \mathcal{U}((0, 1))$, we have $F^{-1}(U) \stackrel{d}{=} X$ for the left-inverse F^{-1} of F

- we don't know F , but are only given samples $X_1, \dots, X_n \stackrel{d}{=} X$
 - naïve approach: replace F by empirical cdf $\hat{F}(x) = \frac{\#\{X_i : X_i \leq x\}}{n}$ and set $\hat{X} = \hat{F}^{-1}(U)$ for an independent $U \sim \mathcal{U}((0, 1))$
 - if $X_{(1)}, \dots, X_{(n)}$ is an increasing ordering of the data set and $U \in [k/n, (k+1)/n]$, then $\hat{X} = X_{(k)}$
- ⇒ algorithm learns the empirical distribution $\mathbb{P}_n = \frac{1}{n} \sum_{i=1}^n \delta_{X_i}$ ⇒ overfitting/no creativity”

To evaluate the performance of an algorithm that outputs $\hat{X} = \hat{T}(U)$, for some $\hat{T} \in \sigma(X_1, \dots, X_n)$ and independent noise U we can

- analyse the rate of convergence of

$$\mathbb{E}\left[d(\hat{T}_\# \mathbb{P}_U, \mathbb{P}_X)\right], \quad d \text{ some probability distance or divergence}$$

- study distance of generated distribution to empirical distribution \mathbb{P}_n
- inspect samples visually

Generative Adversarial Networks (GANs)

Key idea: GANs use a game-theoretic setup:

- **Generator G :** maps random, easy to-sample-from noise U to data space to generate samples
- **Discriminator D :** distinguishes between real samples X and generated samples $G(U)$

Objective: A minimax game is played:

$$\inf_{G \in \mathcal{G}} \sup_{D \in \mathcal{D}} |\mathbb{E}[D(X)] - \mathbb{E}[D(G(U))]|$$

Generative Adversarial Networks (GANs)

Key idea: GANs use a game-theoretic setup:

- **Generator G :** maps random, easy to-sample-from noise U to data space to generate samples
- **Discriminator D :** distinguishes between real samples X and generated samples $G(U)$

Objective: A minimax game is played:

$$\inf_{G \in \mathcal{G}} \sup_{D \in \mathcal{D}} |\mathbb{E}[D(X)] - \mathbb{E}[D(G(U))]|$$

- for $\mathcal{D} = \{D : D \text{ is 1-Lipschitz}\}$ this is equivalent to

$$\inf_{G \in \mathcal{G}} W_1(\mathbb{P}_X, G_\# \mathbb{P}_U) = \inf_{G \in \mathcal{G}} \inf_{\gamma \in \Gamma(\mathbb{P}_X, G_\# \mathbb{P}_U)} \int \|x - y\| \gamma(dx, dy), \quad (\text{Wasserstein-GAN})$$

where $\gamma \in \Gamma(\mathbb{P}_X, G_\# \mathbb{P}_U)$ iff $\gamma(dx \times \mathbb{R}^d) = \mathbb{P}_X(dx)$ and $\gamma(\mathbb{R}^d \times dy) = \mathbb{P}(G(U) \in dy)$

Generative Adversarial Networks (GANs)

Key idea: GANs use a game-theoretic setup:

- **Generator G :** maps random, easy to-sample-from noise U to data space to generate samples
- **Discriminator D :** distinguishes between real samples X and generated samples $G(U)$

Objective: A minimax game is played:

$$\inf_{G \in \mathcal{G}} \sup_{D \in \mathcal{D}} |\mathbb{E}[D(X)] - \mathbb{E}[D(G(U))]|$$

- for $\mathcal{D} = \{D : D \text{ is 1-Lipschitz}\}$ this is equivalent to

$$\inf_{G \in \mathcal{G}} W_1(\mathbb{P}_X, G_\# \mathbb{P}_U) = \inf_{G \in \mathcal{G}} \inf_{\gamma \in \Gamma(\mathbb{P}_X, G_\# \mathbb{P}_U)} \int \|x - y\| \gamma(dx, dy), \quad (\text{Wasserstein-GAN})$$

where $\gamma \in \Gamma(\mathbb{P}_X, G_\# \mathbb{P}_U)$ iff $\gamma(dx \times \mathbb{R}^d) = \mathbb{P}_X(dx)$ and $\gamma(\mathbb{R}^d \times dy) = \mathbb{P}(G(U) \in dy)$

- in practice, \mathcal{D} and \mathcal{G} are chosen as **parameterised neural network classes** and **expectations** are replaced by empirical means:

$$\text{ERM: } \hat{G} \in \arg \min_{G \in \mathcal{G}} \underbrace{\sup_{D \in \mathcal{D}} |\mathbb{E}_n D - \mathbb{E}[D(G(U))]|}_{\approx W_1(\mathbb{P}_n, G_\# \mathbb{P}_U)}$$

Langevin diffusion models

Langevin MCMC algorithm: given target density p_0 simulate diffusion

$$dZ_t = \nabla \log p_0(Z_t) dt + \sqrt{2} dW_t$$

and output Z_T for T “sufficiently large”: if p_0 is sufficiently nice, then $Z_t \xrightarrow{d} p_0$

Langevin diffusion models

Langevin MCMC algorithm: given target density p_0 simulate diffusion

$$dZ_t = \nabla \log p_0(Z_t) dt + \sqrt{2} dW_t$$

and output Z_T for T “sufficiently large”: if p_0 is sufficiently nice, then $Z_t \xrightarrow{d} p_0$

- bad idea: estimate p_0 by some \hat{p} and replace $\nabla \log p_0$ by $\nabla \log \hat{p}$ ([explicit approach](#))

Langevin diffusion models

Langevin MCMC algorithm: given target density p_0 simulate diffusion

$$dZ_t = \nabla \log p_0(Z_t) dt + \sqrt{2} dW_t$$

and output Z_T for T “sufficiently large”: if p_0 is sufficiently nice, then $Z_t \xrightarrow{d} p_0$

- bad idea: estimate p_0 by some \hat{p} and replace $\nabla \log p_0$ by $\nabla \log \hat{p}$ (**explicit approach**)
- **implicit approach:** consider $p_{0,\sigma} = p_0 * \phi_{0,\sigma^2}$ for a Gaussian density ϕ_{0,σ^2} (σ^2 small) instead and target **score** $\nabla \log p_{0,\sigma^2}$ directly

Langevin diffusion models

Langevin MCMC algorithm: given target density p_0 simulate diffusion

$$dZ_t = \nabla \log p_0(Z_t) dt + \sqrt{2} dW_t$$

and output Z_T for T “sufficiently large”: if p_0 is sufficiently nice, then $Z_t \xrightarrow{d} p_0$

- bad idea: estimate p_0 by some \hat{p} and replace $\nabla \log p_0$ by $\nabla \log \hat{p}$ (**explicit approach**)
- **implicit approach:** consider $p_{0,\sigma^2} = p_0 * \phi_{0,\sigma^2}$ for a Gaussian density ϕ_{0,σ^2} (σ^2 small) instead and target score $\nabla \log p_{0,\sigma^2}$ directly
- **denoising score matching:** let $X \sim p_0$, $X_\sigma^2 = X + \sigma \varepsilon \sim p_0 * \phi_{0,\sigma^2}$, for indep. noise $\varepsilon \sim \mathcal{N}(0, \mathbb{I}_d)$

$$\begin{aligned}\nabla \log p_{0,\sigma^2}(x) &= \frac{\int \nabla_x \phi_{0,\sigma^2}(x-y) p_0(dy)}{p_{0,\sigma^2}(x)} = \int \nabla_x \log \phi_{0,\sigma^2}(x-y) \underbrace{\frac{\phi_{0,\sigma^2}(x-y) p_0(dy)}{p_{0,\sigma^2}(x)}}_{=\mathbb{P}(X \in dy | X_\sigma^2 = x)} \\ &= \mathbb{E}[\nabla \log \phi_{X,\sigma^2}(X_\sigma^2) | X_\sigma^2 = x]\end{aligned}$$

Langevin diffusion models

Langevin MCMC algorithm: given target density p_0 simulate diffusion

$$dZ_t = \nabla \log p_0(Z_t) dt + \sqrt{2} dW_t$$

and output Z_T for T “sufficiently large”: if p_0 is sufficiently nice, then $Z_t \xrightarrow{d} p_0$

- bad idea: estimate p_0 by some \hat{p} and replace $\nabla \log p_0$ by $\nabla \log \hat{p}$ (**explicit approach**)
- **implicit approach:** consider $p_{0,\sigma^2} = p_0 * \phi_{0,\sigma^2}$ for a Gaussian density ϕ_{0,σ^2} (σ^2 small) instead and target score $\nabla \log p_{0,\sigma^2}$ directly
- **denoising score matching:** let $X \sim p_0$, $X_\sigma^2 = X + \sigma \varepsilon \sim p_0 * \phi_{0,\sigma^2}$, for indep. noise $\varepsilon \sim \mathcal{N}(0, \mathbb{I}_d)$

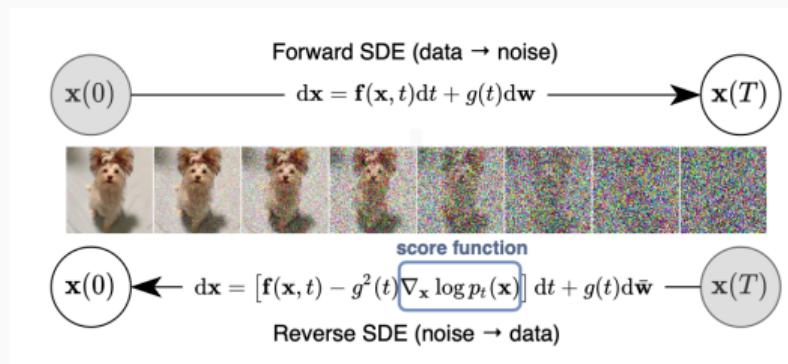
$$\begin{aligned}\nabla \log p_{0,\sigma^2}(x) &= \frac{\int \nabla_x \phi_{0,\sigma^2}(x-y) p_0(dy)}{p_{0,\sigma^2}(x)} = \int \nabla_x \log \phi_{0,\sigma^2}(x-y) \underbrace{\frac{\phi_{0,\sigma^2}(x-y) p_0(dy)}{p_{0,\sigma^2}(x)}}_{=\mathbb{P}(X \in dy | X_\sigma^2 = x)} \\ &= \mathbb{E}[\nabla \log \phi_{X,\sigma^2}(X_\sigma^2) | X_\sigma^2 = x]\end{aligned}$$

$$\text{ERM: } \hat{s} \in \arg \min_{s \in \mathcal{S}} \frac{1}{M} \sum_{j=1}^M \|s(X_{i_j, \sigma^2}) - \nabla \log \phi_{X_{i_j}, \sigma^2}(X_{i_j, \sigma^2})\|^2 = \arg \min_{s \in \mathcal{S}} \frac{1}{M} \sum_{j=1}^M \|s(X_{i_j, \sigma^2}) + \frac{1}{\sigma^2} \varepsilon_{i_j}\|^2,$$

where the X_{i_j} are uniformly sampled from X_1, \dots, X_n and $X_{i_j, \sigma^2} = X_{i_j} + \sigma \varepsilon_{i_j}$

Denoising diffusion models

- provide an **iterative generative algorithm** to create new samples that approximately match the target distribution p_0 , given a finite number of samples corresponding to an unknown p_0
- general idea: find a **stochastic process** that perturbs p_0 to a new distribution p_T such that
 - 1) p_T or a good approximation thereof is **easy to sample from**, and
 - 2) the perturbation is **reversible** in the sense that we know how to **simulate the time-reversed process**



Source: Song et al. (2021). Score based generative modeling through stochastic differential equations. *ICLR*.

Denoising Diffusion Models

- for some fixed time $T > 0$ consider the forward model

$$dX_t = b(t, X_t) dt + \sigma(t, X_t) dW_t, \quad t \in [0, T], X_0 \sim p_0$$

Denoising Diffusion Models

- for some fixed time $T > 0$ consider the **forward model**

$$dX_t = b(t, X_t) dt + \sigma(t, X_t) dW_t, \quad t \in [0, T], X_0 \sim p_0$$

- letting $p_t(x) = \int p_{0,t}(y, x) p_0(dy)$ be the marginal densities of (X_t) , the **time reversal** $\hat{X}_t = X_{T-t}$ solves

$$d\hat{X}_t = -\bar{b}(T-t, \hat{X}_t) dt + \sigma(T-t, \hat{X}_t) d\bar{W}_t, \quad t \in [0, T], \hat{X}_0 \sim p_T,$$

where

$$\begin{aligned}\bar{b}_i(t, x) &= b_i(t, x) - \frac{1}{p_t(x)} \sum_{j,k=1}^d \frac{\partial}{\partial x_j} (p_t(x) \sigma_{ik}(t, x) \sigma_{jk}(t, x)) \\ &= b_i(t, x) - (\nabla \cdot \Sigma(t, x))_i - (\nabla \log p_t(x))_i, \quad i = 1, \dots, d, \Sigma = \sigma \sigma^\top\end{aligned}$$

Denoising Diffusion Models

- for some fixed time $T > 0$ consider the **forward model**

$$dX_t = b(t, X_t) dt + \sigma(t, X_t) dW_t, \quad t \in [0, T], X_0 \sim p_0$$

- letting $p_t(x) = \int p_{0,t}(y, x) p_0(dy)$ be the marginal densities of (X_t) , the **time reversal** $\tilde{X}_t = X_{T-t}$ solves

$$d\tilde{X}_t = -\bar{b}(T-t, \tilde{X}_t) dt + \sigma(T-t, \tilde{X}_t) d\bar{W}_t, \quad t \in [0, T], \tilde{X}_0 \sim p_T,$$

where

$$\begin{aligned}\bar{b}_i(t, x) &= b_i(t, x) - \frac{1}{p_t(x)} \sum_{j,k=1}^d \frac{\partial}{\partial x_j} (p_t(x) \sigma_{ik}(t, x) \sigma_{jk}(t, x)) \\ &= b_i(t, x) - (\nabla \cdot \Sigma(t, x))_i - (\nabla \log p_t(x))_i, \quad i = 1, \dots, d, \Sigma = \sigma \sigma^\top\end{aligned}$$

- time-reversed process solves a **time-inhomogeneous SDE**, now with drift $-\bar{b}(T-t, \cdot)$ involving the **score** $\nabla \log p_t$, which depends on the **unknown** data distribution p_0
- score needs to be estimated from the data

Denoising score matching

- denoising score matching:

$$\begin{aligned}\nabla \log p_t(x) &= \frac{\int \nabla_x p_{0,t}(y, x) p_0(dy)}{p_t(x)} = \int \nabla_x \log p_{0,t}(y, x) \underbrace{\frac{p_{0,t}(y, x) p_0(dy)}{p_t(x)}}_{=\mathbb{P}(X_0 \in dy | X_t = x)} \\ &= \mathbb{E}[\nabla_2 \log p_{0,t}(X_0, X_t) | X_t = x]\end{aligned}$$

and thus

$$\hat{s} := \nabla \log p_t \in \arg \min_{s \text{ meas.}} \mathbb{E}[\|s(X_t) - \nabla_2 \log p_{0,t}(X_0, X_t)\|^2]$$

≈ given data $(X_0^i)_{i \in [n]} \stackrel{\text{iid}}{\sim} p_0$ define the **denoising score estimator**

$$\hat{s} \in \arg \min_{s \in \mathcal{S}} \frac{1}{n} \sum_{i=1}^n \mathbb{E}_{X_0^i} \left[\int_{\underline{T}}^T \|s(t, X_t) - \nabla_2 \log p_{0,t}(X_0, X_t)\|^2 dt \right],$$

where $0 < \underline{T} \ll T$ and \mathcal{S} is an approximating function class, e.g. **space-time neural networks**

Generative process

On $[0, T - \underline{T}]$, simulate

$$dY_t = (-b(T-t, Y_t) + \nabla \cdot \Sigma(T-t, Y_t) + \Sigma(T-t, Y_t) \hat{\mathfrak{s}}(T-t, Y_t)) dt + \sigma(T-t, Y_t) dW_t, \quad \mathbb{P}^{Y_0}(dy) \approx p_T(y) dy$$

Output:

$$Y_{T-\underline{T}} \stackrel{d}{\approx} \hat{X}_{T-\underline{T}} = X_{\underline{T}} \stackrel{d}{\approx} X_0$$

Generative process

On $[0, T - \underline{T}]$, simulate

$$dY_t = (-b(T-t, Y_t) + \nabla \cdot \Sigma(T-t, Y_t) + \Sigma(T-t, Y_t) \hat{\mathfrak{g}}(T-t, Y_t)) dt + \sigma(T-t, Y_t) dW_t, \quad \mathbb{P}^{Y_0}(dy) \approx p_T(y) dy$$

Output:

$$Y_{T-\underline{T}} \xrightarrow{d} \bar{X}_{T-\underline{T}} = X_{\underline{T}} \xrightarrow{d} X_0$$

Minimax optimality of diffusion models

Assumptions on data distribution p_0 with support \mathcal{M} :

- $\text{Leb}(\mathcal{M}) > 0$, \mathcal{M} bounded, $p_0|_{\mathcal{M}} \geq m > 0$ and β -smooth: Oko, Akiyama, Suzuki (*ICML* '23), Dou, Kotekal, Xu and Zhou ('24+) [$d = 1$, no log-factors], Holk, Strauch, LT ('25+) [reflected models]
- $d = 1$, $\mathcal{M} = \mathbb{R}$, p_0 not lower bounded: Zhang et al. ('25, *ICML*)
- \mathcal{M} bounded and \subset linear subspace: Oko, Akiyama, Suzuki (*ICML* '23), Chen et al. (*ICML* '23)
- \mathcal{M} is a d^* -dimensional submanifold: Tang and Yang (*AISTATS* '24), Azangulov, Delegiannidis and Rousseau ('24+) [rates adapt to intrinsic dimension d^*]
- $\log p_0(x) = \sum_{J \subset [d], |J| \leq d^*} f_J(x_J)$, for f_J β -Hölder: Kwon et. al ('25+), Fan, Gu and Li ('25+)
- ... [?]

Generative process

On $[0, T - \underline{T}]$, simulate

$$dY_t = (-b(T-t, Y_t) + \nabla \cdot \Sigma(T-t, Y_t) + \Sigma(T-t, Y_t) \hat{\mathfrak{s}}(T-t, Y_t)) dt + \sigma(T-t, Y_t) dW_t, \quad \mathbb{P}^{Y_0}(dy) \approx p_T(y) dy$$

Output:

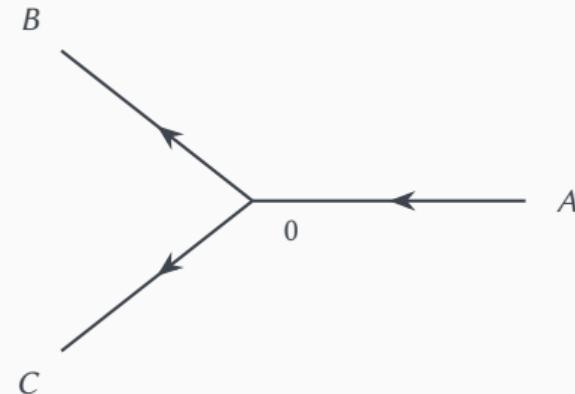
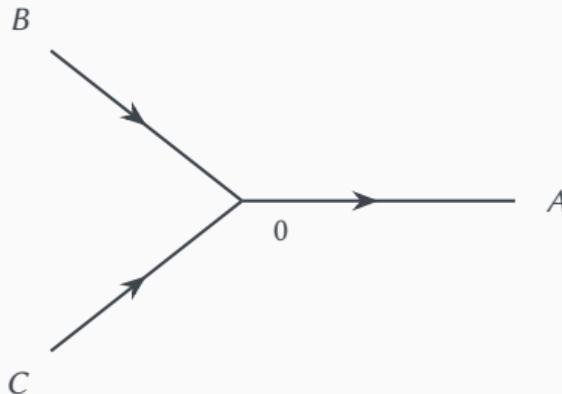
$$Y_{T-\underline{T}} \xrightarrow{d} \hat{X}_{T-\underline{T}} = X_{\underline{T}} \xrightarrow{d} X_0$$

Some fundamental observations

- time reversal at **deterministic** time T forces the backward process to be time-inhomogeneous
- if p_0 has **low-dimensional support** \mathcal{M} , for small t and x close to \mathcal{M} , $\nabla \log p_t(x)$ is approximately **orthogonal** to \mathcal{M} ([Stanczuk et al., ICML '24](#))
- initialising the generative process in a distribution that is not close to \mathbb{P}^{X_T} and simulating for $T - \underline{T}$ time units will not give useful results \rightsquigarrow algorithm is **not adaptive** to the noise level in the data

Homogeneous time reversal

- Markov property: “the past and future of a Markov process are conditionally independent given the present” \rightsquigarrow time-reversed Markov processes are Markov
- to ensure that a homogeneous Markov process remains homogeneous under time reversal, we need to reverse at a suitable random (life)time ζ . This can be
 - a randomised stopping time such as an independent exponential time;
 - a last exit time;
 - a first hitting time;
 - any terminal time, that is, any stopping time T such that $T = t + T \circ \theta_t$ on $\{T > t\}$
- retaining the strong Markov property under time reversal is a bit more tricky:



h -transforms and time reversal

h -transform

For a possibly killed, homogeneous strong Markov process X with state space S , let h be an excessive function, that is

$$\mathbb{E}_x[h(X_t)] \leq h(x) \quad \text{and} \quad \lim_{t \rightarrow 0} \mathbb{E}_x[h(X_t)] = h(x).$$

Then,

$$P_t^h f(x) = \mathbb{E}_x \left[\frac{h(X_t)}{h(x)} f(X_t) \mathbf{1}_{\{X_t \in S\}} \right] \mathbf{1}_{(0, \infty)}(h(x)), \quad f \in \mathcal{B}_b(\mathbb{R}^d),$$

defines a sub-Markov semigroup. The corresponding Markov process X^h is strong Markov and is called h -transform of X .

- suppose that X is a continuous and **self-dual** Feller process (i.e., its generator satisfies $A = A^*$)
- if X^h has a finite killing time ζ , then the time-reversed process $\overset{\leftarrow}{X}_t^h = X_{\zeta-t}^h$ is **homogeneous, strong Markov** and is a \tilde{h} -transform of X .

h-transforming a killed diffusion

- consider a **symmetric** diffusion process

$$dX_t = b(X_t) dt + \sigma(X_t) dW_t$$

with invariant measure m and let Z be its version **killed at an independent exponential time** with parameter $r > 0$

- as an excessive function for Z use

$$h(x) = \int G_r(x, y) \kappa(dy)$$

for the **Green kernel** $G_r(x, y) = \int_0^\infty e^{-rt} p_t(x, y) dy$ and a **representing measure** κ

- $\kappa(dy) = r dy \rightsquigarrow h = 1$ and $Z^h = Z$
- $\kappa(dy) = \frac{1}{G_r(x_0, y)} \beta(dy) \rightsquigarrow Z$ conditioned to have distribution β before killing if started in x_0
- Z is a killed Brownian motion and $\kappa(dy) = \sigma_R(dy)$ for the surface measure σ_R of an R -sphere $\mathbb{S}^{d-1}(R) \rightsquigarrow Z^h$ is killed at last exit from $\mathbb{S}^{d-1}(R)$

A time-homogeneous generative process

Proposition (Christensen, Strauch and LT (2025+))

1. Z^h is an Itô-diffusion with dynamics

$$dZ_t^h = (b(Z_t^h) + \Sigma(X_t) \nabla \log h(X_t)) dt + \sigma(Z_t^h) dW_t$$

outside $\text{supp } \kappa$ and its distribution at the lifetime is given by

$$\mathbb{P}_x(Z_{\zeta^-}^h \in dy) = \frac{G_r(x, y)}{h(x)} \kappa(dy)$$

2. Let $\alpha = \mathbb{P}^{Z_0^h}$. Then Z_t^h is an \hat{h} -transform of Z with initial distribution $\mathbb{P}_\alpha(Z_{\zeta^-}^h \in dy)$ and

$$\hat{h}(x) := \int \frac{G_r(x, y)}{h(y)} \alpha(dy).$$

In particular, Z_t^h has dynamics

$$dZ_t^h = (b(Z_t^h) + \Sigma(Z_t^h) \nabla \log \hat{h}(Z_t^h)) dt + \sigma(Z_t^h) d\bar{W}_t,$$

outside $\text{supp } \alpha =: \mathcal{M}$ and $\mathbb{P}_\alpha(Z_{\zeta^-}^h \in dy \mid Z_0^h = x) = \frac{G_r(x, y)}{\hat{h}(x)h(y)} \alpha(dy)$ for $\mathbb{P}_\alpha(Z_{\zeta^-}^h \in \cdot)$ -a.e. x .

A time-homogeneous generative process

Idealised algorithm:

1. Initialise $Z_0^{\tilde{h}} \sim \tilde{\beta} \approx \mathbb{P}_\alpha(Z_{\zeta^-}^h)$
 - for ergodic forward process with stationary distribution μ and small exponential killing rate $r > 0$, choose $\tilde{\beta} = \mu$ [↔ ergodic diffusion model]
 - for exponentially killed BM with small killing rate $r > 0$, choose $\tilde{\beta} = \text{Laplace}(0, (2r)^{-1/2} I_d)$ [↔ variance exploding diffusion model]
 - for $\kappa(dy) = \frac{1}{G_r(x_0, y)} \delta_z$, choose $\tilde{\beta} = \delta_z$
2. Simulate diffusion $Z^{\tilde{h}}$ until killing time and output $Z_{\zeta^-}^{\tilde{h}}$

A time-homogeneous generative process

Idealised algorithm:

1. Initialise $Z_0^{\tilde{h}} \sim \tilde{\beta} \approx \mathbb{P}_\alpha(Z_{\zeta^-}^h)$
 - for ergodic forward process with stationary distribution μ and small exponential killing rate $r > 0$, choose $\tilde{\beta} = \mu$ [↔ ergodic diffusion model]
 - for exponentially killed BM with small killing rate $r > 0$, choose $\tilde{\beta} = \text{Laplace}(0, (2r)^{-1/2} I_d)$ [↔ variance exploding diffusion model]
 - for $\kappa(dy) = \frac{1}{G_r(x_0, y)} \delta_z$, choose $\tilde{\beta} = \delta_z$
2. Simulate diffusion $Z^{\tilde{h}}$ until killing time and output $Z_{\zeta^-}^{\tilde{h}}$

Requirements for implementation

1. learn $\nabla \log \tilde{h}$ (only a function in space – no time component);

A time-homogeneous generative process

Idealised algorithm:

1. Initialise $Z_0^{\tilde{h}} \sim \tilde{\beta} \approx \mathbb{P}_\alpha(Z_{\zeta^-}^h)$
 - for ergodic forward process with stationary distribution μ and small exponential killing rate $r > 0$, choose $\tilde{\beta} = \mu$ [\leftrightarrow ergodic diffusion model]
 - for exponentially killed BM with small killing rate $r > 0$, choose $\tilde{\beta} = \text{Laplace}(0, (2r)^{-1/2} I_d)$ [\leftrightarrow variance exploding diffusion model]
 - for $\kappa(dy) = \frac{1}{G_r(x_0, y)} \delta_z$, choose $\tilde{\beta} = \delta_z$
2. Simulate diffusion $Z^{\tilde{h}}$ until killing time and output $Z_{\zeta^-}^h$

Requirements for implementation

1. learn $\nabla \log \tilde{h}$ (only a function in space – no time component);
2. learn killing time ζ of $Z^{\tilde{h}}$

Learning to kill

Polarity hypothesis

Assume that $\mathcal{M} = \text{supp } \alpha$ is **polar** for X , i.e., for any $x \in \mathbb{R}^d$, $\mathbb{P}_x(\inf\{t > 0 : X_t \in \mathcal{M}\} < \infty) = 0$.

Learning to kill

Polarity hypothesis

Assume that $\mathcal{M} = \text{supp } \alpha$ is **polar** for X , i.e., for any $x \in \mathbb{R}^d$, $\mathbb{P}_x(\inf\{t > 0 : X_t \in \mathcal{M}\} < \infty) = 0$.

Theorem (Christensen, Strauch and LT (2025+))

Under the polarity hypothesis, the backward process Z^h is **killed at first entrance into \mathcal{M}** .

Learning to kill

Polarity hypothesis

Assume that $\mathcal{M} = \text{supp } \alpha$ is **polar** for X , i.e., for any $x \in \mathbb{R}^d$, $\mathbb{P}_x(\inf\{t > 0 : X_t \in \mathcal{M}\} < \infty) = 0$.

Theorem (Christensen, Strauch and LT (2025+))

Under the polarity hypothesis, the backward process $Z^{\hat{h}}$ is **killed at first entrance into \mathcal{M}** .

Possible strategies to estimate a δ -fattening $\mathcal{M}_\delta = \{x : \text{dist}(x, \mathcal{M}) \leq \delta\}$ given data $X^1, \dots, X^n \stackrel{\text{iid}}{\sim} \alpha$ and an estimator \hat{s} of $s := \nabla \log \hat{h}$:

- **explicit plug-in approach:** estimate \mathcal{M}_δ directly or indirectly by setting $\widehat{\mathcal{M}}_\delta = (\widehat{\mathcal{M}})_\delta$; then set $\hat{\zeta} := \inf\{t \geq 0 : Z_t^{\hat{s}} \in \widehat{\mathcal{M}}_\delta\}$
- **implicit approach:** use explosive behaviour of s as $x \rightarrow \mathcal{M}$

Theorem (Christensen, Kallsen, Strauch and LT (2025+))

Suppose that \mathcal{M} is polar for X and Y solving $dY_t = \sigma(Y_t) dB_t$. Then, it a.s. holds that

$$\zeta = \inf \left\{ t \geq 0 : \sup_{s \leq t} |\hat{s}(Z_s^{\hat{h}})| = \infty \right\} = \inf \left\{ t \geq 0 : \|\hat{s}(Z^{\hat{h}})\|_{L^2([0,t])} = \infty \right\}.$$

Denoising score matching

- for $\mathbb{P}_\alpha(Z_{\zeta^-}^h \in \cdot)$ -a.e. x

$$\begin{aligned}\mathfrak{s}(x) &= \nabla \log \tilde{h}(x) = \frac{1}{\tilde{h}(x)} \int \nabla_x G_r(x, y) \frac{1}{h(y)} \alpha(dy) = \int \nabla_x \log G_r(x, y) \frac{G_r(x, y)}{\tilde{h}(x) h(y)} \alpha(dy) \\ &= \mathbb{E}[\nabla_x \log G_r(x, Z_{\zeta^-}^h) | Z_0^h = x] \\ &= \mathbb{E}_\alpha[\nabla_x \log G_r(x, Z_0^h) | Z_{\zeta^-}^h = x]\end{aligned}$$

- this implies that on $\mathbb{R}^d \setminus \mathcal{M}_\delta$, \mathfrak{s} agrees $\mathbb{P}_\alpha(Z_{\zeta^-}^h \in \cdot)$ -a.e. with the minimiser of

$$\mathcal{B}(\mathbb{R}^d; \mathbb{R}^d) \ni s \mapsto \mathbb{E}_\alpha \left[\|s(Z_{\zeta^-}^h) - \nabla \log G_r(Z_0^h, Z_{\zeta^-}^h)\|^2 \mathbf{1}_{\{\|Z_{\zeta^-}^h - Z_0^h\| > \delta\}} \right]$$

- note that if $Z^h = Z$, then $\zeta \sim \text{Exp}(r)$ independent of X , $Z_{\zeta^-} = X_\zeta$ has full support and we have

$$\mathbb{E}_\alpha \left[\|s(Z_{\zeta^-}^h) - \nabla \log G_r(Z_0^h, Z_{\zeta^-}^h)\|^2 \mathbf{1}_{\{\|Z_{\zeta^-}^h - Z_0^h\| > \delta\}} \right] = r \mathbb{E}_\alpha \left[\int_0^\zeta \|s(Z_t^h) - \nabla \log G_r(Z_0^h, Z_t^h)\|^2 \mathbf{1}_{\{\|Z_t^h - Z_0^h\| > \delta\}} dt \right]$$

Projection learning

- we don't have to start the backward process approximately in $\mathbb{P}_\alpha(Z_\zeta^h \in dy)$: it will always be killed on the data support \mathcal{M} and different initialisations will yield different output distributions supported on $\mathcal{M} \rightsquigarrow$ **natural conditioning**
- a natural question is therefore what happens if we don't start the generative process from pure noise but something more informative, say a **masked** or **moderately noised** picture



- it turns out that the natural conditioning aspect entails a **blessing of dimensionality**

Projection learning

Let Z be an exponentially killed Brownian motion. Then,

$$\tilde{h}(x) = \int G_r(x, y) \alpha(dy), \quad G_r(x, y) = 2(2\pi)^{-d/2} r \left(\frac{\sqrt{2r}}{|x-y|} \right)^{\frac{d-2}{2}} K_{\frac{d-2}{2}} \left(\frac{\sqrt{2r}}{|x-y|} \right).$$

For large d ,

$$\nabla \log \tilde{h}(x) \approx d \frac{\int \frac{x-y}{|x-y|^d} \alpha(dy)}{\int |x-y|^{2-d} \alpha(dy)}$$

and thus, if there is a unique projection $x^* \in \arg \min_{y \in \mathcal{M}} |x - y|$ of x onto \mathcal{M} , then

$$\nabla \log \tilde{h}(x) \approx d \frac{x^* - x}{|x^* - x|^2} = d \frac{\text{sign}(x^* - x)}{|x^* - x|}$$

Theorem (Christensen, Kallsen, Strauch and LT (2025+))

Let $\delta, \varepsilon > 0$ and fix an observation $x \in \mathbb{R}^d$. If $\alpha(B(x, r)) > \varepsilon$ for some ball $B(x, r)$ with radius $r > 0$ around y , then

$$\mathbb{P}\left(Z_{\zeta^-}^{\tilde{h}} \in \mathcal{M} \cap B(x, (1 + \delta)r) \mid Z_0^{\tilde{h}} = x\right) \geq 1 - \frac{1}{\varepsilon} (1 + \delta)^{2-d}.$$

Projection learning

Consider now estimators \hat{s}_n , an independent Brownian motion W and let $\hat{Z}^{\hat{s}_n}$ be the process solving

$$d\hat{Z}_t^{\hat{s}_n} = \hat{s}_n(\hat{Z}_t^{\hat{s}_n}) \mathbf{1}_{\{t \leq \hat{\zeta}\}} dt + \mathbf{1}_{\{t \leq \tilde{\zeta}\}} dW_t, \quad \hat{\zeta} := \inf \left\{ t \geq 0 : \|\hat{Z}^{\hat{s}_n}\|_{L^2[0,t]} > M \right\}.$$

Theorem (Christensen, Kallsen, Strauch and LT (2025+))

Fix an observation $x \in \mathbb{R}^d$. Suppose that

- for any $\tilde{\delta}, \delta, \varepsilon > 0$ it holds for sufficiently large n that

$$\mathbb{P} \left(\left\| (\hat{s}_n(Z^{\tilde{h}}) - s(Z^{\tilde{h}})) \mathbf{1}_{\{Z^{\tilde{h}} \notin \mathcal{M}_{\tilde{\delta}}\}} \right\|_{L^2(\zeta)} > \delta \mid Z_0^{\tilde{h}} = x \right) < \varepsilon$$

- for any $n \in \mathbb{N}$ and $\tilde{\delta} > 0$, the function \hat{s}_n is $L_{\tilde{\delta}}$ -Lipschitz on $\mathcal{M}_{\tilde{\delta}}^c$

Let $\delta, \varepsilon, \tilde{\delta}, \tilde{\varepsilon} > 0$. If $\alpha(B(x, r)) > \varepsilon$, then, for sufficiently large $M > 0$ and $n \in \mathbb{N}$,

$$\mathbb{P}(\hat{Z}_{\hat{\zeta}}^{\hat{s}_n} \in \mathcal{M}_{\tilde{\delta}} \cap B(x, (1 + \delta)r) \mid \hat{Z}_0^{\hat{s}_n} = x) > 1 - \frac{1}{\varepsilon}(1 + \delta)^{2-d} - \tilde{\varepsilon}.$$

Projection learning

Consider now estimators \hat{s}_n , an independent Brownian motion W and let $\hat{Z}^{\hat{s}_n}$ be the process solving

$$d\hat{Z}_t^{\hat{s}_n} = \hat{s}_n(\hat{Z}_t^{\hat{s}_n}) \mathbf{1}_{\{t \leq \hat{\zeta}\}} dt + \mathbf{1}_{\{t \leq \tilde{\zeta}\}} dW_t, \quad \hat{\zeta} := \inf \left\{ t \geq 0 : \|\hat{Z}^{\hat{s}_n}\|_{L^2[0,t]} > M \right\}.$$

Theorem (Christensen, Kallsen, Strauch and LT (2025+))

Fix an observation $x \in \mathbb{R}^d$. Suppose that

- for any $\tilde{\delta}, \delta, \varepsilon > 0$ it holds for sufficiently large n that

$$\mathbb{P} \left(\left\| (\hat{s}_n(Z^{\tilde{h}}) - s(Z^{\tilde{h}})) \mathbf{1}_{\{Z^{\tilde{h}} \notin \mathcal{M}_{\tilde{\delta}}\}} \right\|_{L^2(\zeta)} > \delta \mid Z_0^{\tilde{h}} = x \right) < \varepsilon$$

- for any $n \in \mathbb{N}$ and $\tilde{\delta} > 0$, the function \hat{s}_n is $L_{\tilde{\delta}}$ -Lipschitz on $\mathcal{M}_{\tilde{\delta}}^c$

Let $\delta, \varepsilon, \tilde{\delta}, \tilde{\varepsilon} > 0$. If $\alpha(B(x, r)) > \varepsilon$, then, for sufficiently large $M > 0$ and $n \in \mathbb{N}$,

$$\mathbb{P}(\hat{Z}_{\hat{\zeta}}^{\hat{s}_n} \in \mathcal{M}_{\tilde{\delta}} \cap B(x, (1 + \delta)r) \mid \hat{Z}_0^{\hat{s}_n} = x) > 1 - \frac{1}{\varepsilon}(1 + \delta)^{2-d} - \tilde{\varepsilon}.$$

Thank you for your attention!